

FROM A POLISHED PROTOTYPE TO A FAILURE-REVIEW CANDIDATE

Why the iteration changed

The first 72-hour build proved a usable shared-care workflow. Real-clinic review then asked a harder question: when the system is wrong, stale, contradictory, or unavailable, can a care team see the boundary and recover safely?

Decision: narrow, inspectable safety slices

Rounds 1–4 selected explicit allergy contradiction, exposure visibility, provider failure, patient publication, and immutable provenance. Round 9 repaired the production logging boundary before the existing service was updated.

BASELINE**16**

adversarial real-clinic scenarios were audited.

CHOSEN SHAPE**5**

bounded vertical slices connect source, review, and privacy.

SAFETY POSTURE**Human**

AI remains a suggestion; clinical truth stays reviewable.

What the review exposed

- 1 Text can contradict itself.**
A bounded penicillin assertion slice makes both immutable sources visible and gives only a Clinician the adjudication boundary.
- 2 Attention can be mistaken for risk.**
Glance priority, explicit risk, protected attention, and exposure impressions remain separate fields; the view is capped at six.
- 3 Accept is not publish.**
The patient-facing dosage path has a typed draft, clinician approval, explicit publish, recall, correction, and safe projection.
- 4 A provider can fail.**
Redaction precedes an optional provider call; bounded retries and a circuit state surface failure without silently returning fixture output.
- 5 Derived text must remain checkable.**
Every highlight, assertion, and archival summary points to an immutable source version and exact Unicode-codepoint span.


The iteration is successful when a reviewer can answer “what happened, what is still true, and what must a human decide?” without trusting a decorative badge.

This brief reports synthetic prototype evidence. It is not a clinical validation, a general clinical NLP claim, or a production compliance certification.

ONE SOURCE GRAPH, SEVERAL REVIEW GATES

Architecture that keeps trust local

The product does not turn a summary, rank, suggestion, or publication draft into an authoritative source. It keeps the source graph immutable and makes every downstream transition explicit.



Round 9 release evidence

DATABASE GATE PostgreSQL 18 · exact source CI 33650978171 c6e9851288c745ceb66dad32078d1385ffbe3424 Fresh/legacy migration, seed idempotency, schema checks, bridge probe, backend and static gates passed.	HOSTED DEPLOYMENT Existing Render service · Live deploy dep-dac4dgek1f9s73e4qu30 main c6e9851 · rc5 HTTPS redirect, anonymous boundaries, role canaries, and 15/15 sustained asset watch passed.
APPLICATION QUALITY 179 backend tests Focused logging regression: 12 passed. Frontend: 45 Vitest; browser: 18 core + 4 Voice + 2 publication. Global app coverage is 83.30%; runtime application slice excluding standalone scripts is 92.9%.	LOCAL PERFORMANCE BOUNDARY P95 56.053 ms Prior local SQLite/Uvicorn warm-path evidence, 1,000 requests at concurrency 10, zero errors. The authenticated hosted benchmark is pending; this is not a hosted PostgreSQL SLA.

Logging recovery

The incident was not solved by silencing access logs. RED commit f72593c reproduced Uvicorn AccessFormatter failures after a sanitizer flattened `LogRecord.args`. GREEN commit 43714a5 preserves the formatter-compatible argument shape, sanitizes query values separately, and uses a formatter-safe fail-closed fallback.

No migration, dependency manifest, requirements brief, Voice path, DeepSeek configuration, or data-model meaning changed in this repair.

WHAT SURVIVES, WHAT REMAINS PARTIAL

Useful evidence, bounded claims

The system is strong where the contract is narrow and observable. The remaining gaps are named here so a reviewer can distinguish a safety slice from a clinical promise.

CASE	REAL-CLINIC QUESTION	STATUS	CURRENT INTERPRETATION
01–03	Identity, tenant boundary, PHI beyond redaction	PARTIAL	Seeded email identity and app-level scope are tested; no full no-email onboarding, RLS, or third-party retention audit.
04	Redaction ordering	SURVIVES	Typed redaction precedes optional provider calls; failure makes zero provider calls.
05–06	Clinic onboarding, multilingual consult	DOES NOT	Deferred identity onboarding and multilingual ASR/capture are not implemented.
07	Allergy contradiction	PARTIAL	One closed-vocabulary penicillin slice with two sources and clinician review; not general semantic NLP.
08–09	Model hang, provider outage	PARTIAL	Bounded timeout/circuit/degraded UI; no durable queue, replay, or async worker guarantee.
10	Concurrent editing	SURVIVES	Same-section stale writes return deterministic 409 and preserve both submissions.
11	Appointment link delivery	DOES NOT	No delivery, receipt, retry, or acknowledgement channel.
12–15	Dosage, contradiction, priority, exposure bias	PARTIAL	Typed publication, bounded assertions, safety floor, and metadata impressions exist; no general medication NLP, calibration, or debiasing claim.
16	Edited source provenance	SURVIVES	Immutable version and exact codepoint span remain resolvable after edits.

ROUND 9 PARTIAL GATES

Do not overstate live evidence

- The hosted authenticated Glance/Patient benchmark is pending because no permitted authenticated request channel was available.
- Accumulated synthetic rehearsal state hid protected conflict/publication actions from the live top-six; no production cleanup was performed to manufacture them.
- `real-clinic-live1` was not created.

NEXT VALIDATION PRIORITIES

Close the evidence, not the caveat

- Run a bounded authenticated hosted read benchmark through an approved non-secret mechanism.
- Rehearse the protected production state only with an explicit, auditable data-state decision.
- Complete final artifact audit and packaging in Round 10.

Not claimed: live LLM quality, live ASR, microphone capture, FHIR conformance, external delivery receipts, clinical validation, or production compliance.